

SAMPLE REPORT • ANONYMIZED

See what an attacker already knows — before the strike.

Below is a **real, authorized HYDRA engagement** with the client fully anonymized — same depth, same discipline, zero PII. This is exactly what lands in your inbox when you work with us.

Passive only

Receipts-first

GDPR • NDA-ready

2-3 days

⚠ **SAMPLE • ANONYMIZED • ILLUSTRATIVE.** This is a real HYDRA engagement with the client identity, domains, IP addresses and personnel **fully redacted**. The structure, methodology, lane logic and finding **types** are unchanged — this is exactly what a HYDRA report delivers. No live target is identified.

TIER 3 • FULL INTELLIGENCE • PASSIVE EXTERNAL EXPOSURE

HYDRA
by TIA

Sample Engagement Report

Authorized passive external exposure analysis — 5-lane HYDRA methodology, evidence-first.

Method: **Passive OSINT** Lanes: **5** Findings: **12** Basis: **real engagement, de-identified**

ENGAGEMENT • SCOPE

Client: [REDACTED] — a mid-size Prague-based multi-branch consumer group (~5 owned domains, ~40 staff, a routine supplier-payment workflow). Identity, domains, IPs and people redacted for this sample.

Authorization: written owner consent (assessment of the client's own public exposure). **Method:** 5-lane HYDRA — **PASSIVE only** (DNS/CT/DoH, public web, response headers, RDAP/whois, OSINT). No active scanning, no authentication attempts, no exploitation. Receipts-first, zero fabrication.

Declared up front: account-level breach confirmation and active characterization of the on-prem asset are **not** possible passively → scoped to a paid, signed Phase 2.

EXTERNAL THREAT EXPOSURE

ELEVATED

Four critical exposures on an otherwise healthy perimeter. The realistic money-loss path — **email spoofing** → **fake supplier invoice** — was confirmed independently across three lanes.

12

FINDINGS

4

CRITICAL / HIGH

5

RECON LANES

Passive

METHOD — 0 TOUCH

EXECUTIVE SUMMARY — 4 REAL RISKS

The rest of the perimeter is reassuringly healthy (see the honest-negatives section). Four exposures carry real, exploitable risk:

- 1 On-prem remote-access box exposed to the internet.** A management endpoint on consumer broadband, end-of-life web server, self-signed certificate — the single most attackable real asset. **CRITICAL**
- 2 Email spoofing wide open on both flagship domains** (DMARC p=none). Anyone can send mail that lands in the inbox appearing to come from the founder or any branch → Business Email Compromise / fake-invoice. *Triple-confirmed across infra + dark-int + people lanes.* **CRITICAL**
- 3 Third-party look-alike domain with live mail.** A recently-registered near-spelling of the brand, behind a privacy proxy, with active mail records — a ready-made phishing channel impersonating the client. **CRITICAL**
- 4 Shared branch mailboxes + no enforced MFA.** The credential-hygiene foothold that makes risk #2 trivial to weaponize. **HIGH**

FINDINGS BY SEVERITY

F1 On-prem remote-access endpoint exposed

[INFRA]

A deliberate DNS record points a management hostname at **consumer broadband (not a datacenter)**. Passive fingerprint: end-of-life web server, self-signed certificate (multi-year). Classic on-prem management / VPN / back-office box on a flat office network likely touching POS, invoicing and payroll — internet-reachable, unadvertised in certificate transparency. → identify it, place behind a zero-trust broker, patch/replace, pull off the public net. *What it actually runs = Phase 2 (signed scope).*

F2 DMARC p=none on both flagship domains → BEC

[INFRA · DARK-INT · PEOPLE]

Both flagship domains publish **p=none** (monitor, no enforcement) despite valid SPF and DKIM. Anyone can send **From: <anyone>@[client-domain]** and land in the inbox. For a multi-branch business paying suppliers, this is a

direct fake-invoice / payment-redirect vector. → add aggregate reporting, confirm sender alignment, then move **none** → **quarantine** → **reject** on both.

F3 Look-alike domain, mail-capable

[BRAND · INFRA · WEB]

A near-spelling of the brand, registered ~6 months ago behind a US privacy proxy, on a different hosting stack than the client — every fingerprint diverges, so it is *not* theirs. Storefront disabled, but **live MX records** present: a disabled store ≠ disabled mail = a ready phishing/BEC channel. → attempt acquisition; else abuse-report + registry dispute; monitor records; warn staff and customers.

F4 Shared branch mailboxes + MFA hygiene

[DARK-INT · PEOPLE]

Multiple shared role mailboxes confirmed (branch + **info@** / **orders@** / **catering@**). These are the highest-reuse, lowest-rotation accounts — shared passwords, rarely changed on departure, reused as portal logins, seldom per-user MFA. One leaked branch password lets an attacker read supplier invoices and reservation PII → invoice fraud (amplifies F2). → enforce org-wide 2-step verification; migrate shared → owner+delegation; rotate now.

F5 Expired / wrong-owner TLS + HTTPS→HTTP downgrade

[WEB · INFRA]

Two secondary sites present an expired certificate for an unrelated owner and downgrade **https** → **plaintext http**. Erodes trust and trains users to click through certificate warnings. → valid certs or TLS-valid edge + clean redirects.

F6 Shop domain with no SPF / no DMARC

[INFRA]

A secondary shop domain is fully spoofable (no SPF, no DMARC) — a clean channel for customer-facing phishing. → **v=spf1 -all + p=reject**.

F7 Single publicly-named finance identity

[PEOPLE]

Public professional profiles name exactly one finance decision-maker — both an impersonation *target* and a spoofing *vector* for payment fraud. → authenticate finance comms; callback/codeword for any payment or bank-detail change; reduce public identifiability.

F8 Third-party processor cascade

[DARK-INT]

Several third-party processors sit in the CRM/POS/reservations/mail path (guest PII, payment context). A breach at any one leaks client data indirectly; password reuse turns it into a mail-tenant compromise. → inventory staff↔processor logins; unique passwords + MFA; breach-monitor admin accounts.

F9 Content-Security-Policy report-only; weak Referrer-Policy

[WEB · INFRA]

The primary site ships CSP in report-only mode and a permissive Referrer-Policy. → promote CSP to enforced (the sister site already does); tighten Referrer-Policy.

F10 Defensive-registration & listing gaps

[BRAND]

Several obvious brand variants (.com/.eu/.net + alternate spellings) are unregistered and a social handle is unclaimed — open ground for impersonation. → register defensives, claim the handle, verify business listings.

F11 Broken backup MX record

[INFRA]

A secondary MX points to a non-resolving host (possibly tied to the F1 on-prem box). → fix or remove.

F12 Account-level breach exposure — UNVERIFIED

[DARK-INT]

Free-tier sources show **no public hit** (confirmed negative via dorks + public leak lists), but regional combolists sit behind paid databases. This is absence of public evidence, *not* evidence of absence. → Phase 2 paid pass (breach DBs, Split-Brain on any hit).

HONEST NEGATIVES — WHAT'S ACTUALLY SOLID

A report that is all-red is selling fear, not truth. Here's the genuinely good news, evidence-backed:

- ✓ Managed SaaS platforms eliminate whole classes of self-hosted vulnerability; card data sits with the payment provider.
- ✓ No secrets, stack traces or sensitive paths exposed in page source or robots/sitemap.
- ✓ Mail baseline is healthy — SPF + DKIM present on owned domains; two sub-brands already enforce DMARC.
- ✓ No social-media impersonation found; aggregator listings consistent.
- ✓ The owned domain portfolio is genuinely owned and correctly configured.

REMEDIATION PLAN — SEVERITY-ORDERED

● NOW — QUICK WINS

1. Investigate + isolate the on-prem remote box (zero-trust, patch, off the public net).
2. DMARC none → quarantine → reject on both flagship domains (after alignment monitoring).
3. Decide the look-alike domain: acquire, or abuse-report + dispute; warn staff and customers.
4. Enforce org-wide 2-step verification; rotate shared branch passwords.
5. Callback-verification + two-person rule for ANY payment or bank-detail change — the single highest-leverage human control, and it defends even against look-alike-domain attacks that DMARC can't stop.

● SOON

1. Fix expired certs / HTTPS downgrade; add SPF + DMARC to the shop domain.
2. Enforce CSP + tighten Referrer-Policy; register defensive domains + claim listings/handle.
3. Inventory staff↔processor logins; unique passwords + MFA everywhere.

● HYGIENE

1. Remove the broken backup MX; consolidate the domain portfolio under one registrar with 2FA + auto-renew.
2. Verify DNSSEC on remaining domains; confirm consent-gating on tracking.

PHASE 2 — WHERE THIS GOES NEXT (PAID, SIGNED SCOPE)

- Active vulnerability assessment of the on-prem box — what it runs, harden or remove.
- Paid breach-database pass — confirm or deny account-level credential exposure.
- DMARC enforcement rollout + staff BEC-awareness training ("finance never changes a bank account by message").
- Ongoing monitoring — DMARC aggregate reports + brand/domain monitoring.

METHODOLOGY & DISCIPLINE

Receipts-first, zero fabrication. Every / is backed by a passive DNS record, RDAP entry, response header or certificate read. Honest negatives are stated as negatives. Explicitly-unverified items (what the on-prem box runs; account-level breach; true owner of the look-alike domain) are flagged as **unverified** and routed to a signed Phase 2 — never guessed, never inflated.

Passive only. No active exploitation was attempted. All findings derive from publicly available information and open-source intelligence.

HYDRA CUTS HEADS OF DRAGONS — NOT MERCHANTS

We find what an attacker already knows about you, prove it with evidence, and hand you a closed-loop path to fix it. **The first finding is free.**

 HYDRA — INTELLIGENCE CUTS THAT CLOSE CASES • AN INTELLIGENCE PRODUCT BY TIA

SAMPLE / ANONYMIZED. This document reproduces the structure, methodology and finding types of a real, authorized HYDRA engagement; the client identity, domains, IP addresses, certificates and personnel have been fully redacted or generalized, and no live target is identifiable.

Assessment performed using passive reconnaissance only — no active exploitation. hydra.tia-framework.com